Making your tweets more ironic

Kristian Djaković, Rene Kustura, Toma Puljak

University of Zagreb, Faculty of Electrical Engineering and Computing
Unska 3, 10000 Zagreb, Croatia
{author}@fer.hr

Abstract

This document provides the instructions on formatting the TAR system description paper in LaTeX. This is where you write the abstract (i.e., summary) of the work you carried out within the project. The abstract is a paragraph of text ranging between 70 and 150 words.

1. Introduction

Irony is complicated. In general, irony refers to a clash between expectation and outcome of an expression. (Kreuz, 2020). While this definition looks simple enough, detecting irony in text proves to be a difficult task, not only for models, but for people as well. This paper tackles the question: "What makes tweets ironic?" by analysing attention values outputted from a classification model.

The attention mechanism allows modeling of dependencies without regard to their distance in the input or output sequence (Vaswani et al., 2017). What's great about attention is that it adds a layer or explainability to the model. Attention values can be used to interpret the focus of a model when completing various tasks. In this paper, attention values are used to gather most ironic tokens from data samples. Most ironic tokens being those with the highest attention values.

For classification purposes, we used a pretrained RoBERTa Transformer model that was trained on the SemEval2018 Irony Detection dataset (Liu et al., 2019). The model uses multi-head attention, thus we used the *attention rollout* method to aggregate attention values across heads for single tokens and connections between token pairs.

Once we gathered the most ironic tokens, we applied them to tweets that the model classified as *non ironic* and observed the change between the prediction certainty. We use this change in prediction certainty as a metric to evaluate how much a token, or token pair, impacts the ironic sentiment of a given tweet.

2. Background

In this section, we provide an overview of the classification model setup, attention aggregation methods and the datasets evaluate our analysis goal.

2.1. Model Setup

For classification purposes, we used a pretrained RoBERTa Transformer model that had the best performance in irony detection. The model can be found here: `https://huggingface.co/cardiffnlp/twitter-roberta-base-irony`. RoBERTa is a retrained BERT model with improved performance. The details of the architecture can be found in (Liu et al., 2019).

2.2. Attention as an Interpretation Metric

Using attention as an interpretation metric is arguable and research by (Serrano and Smith, 2019) shows that it doesn't necessarily produce correct importance ranking. Nonetheless, we use attention values to interpret how ironic are the tokens, or token pairs.

We considered the alternative of analysing all tokens from selected ironic tweets, but that would require a lot of computation time which we leave for future work.

2.3. Attention Aggregation

To aggregate attention values into a singular output we use the *attention rollout* method. Attention rollout is a recursive method that calculates the product of all attention weights in the attention graph. (Abnar and Zuidema, 2020) This method will compute a singular output for a single token (node) that we feed into it and if we squeeze the resulting matrix without extraction, we get attention values for two connected tokens in the input sentence.

2.4. Data

The dataset used to train the RoBERTa model for irony detection is the SemEval2018 Irony Detection dataset (Hee et al., 2018) which is presplit into training and testing sets. Each tweet is annotated with a binary label of 0 or 1, representing a non-ironic and ironic label respectively.

We applied a simple preprocessing step for all tweets which changes all user tags to *@user* and removes all URLs. This step ensures that tagging specific users in tweets doesn't affect the tokenization and that URLs don't affect the model at all because they, by themselfs, do not add any semantic value to the tweet.

3. Analysing the Most Ironic Tokens

In this section, we will cover the process of finding the most ironic tokens from the used dataset and the methods used for evaluating the results.

3.1. Finding the Tokens

The first step in finding the tokens is to extract all tweets that were annotated as ironic from the training and testing datasets, keeping them separate so we can test the results of the experiment separately. We then classified those tweets using the RoBERTa model to obtain a confidence score for each label. Sorting the tweets regarding to their ironic label

confidence score we obtained a list of the most ironic tweets in the datasets, accoring to the model. The five most ironic tweets from the training and testing datasets can be seen in Table 1 and 2 respectively.

Once we obtained the list of the most ironic tweets, we performe attention aggregation methods outlined in Section 2.3. and sort the tokens by the attention score in descending order. That gives us a list of singular and two connected tokens that the model focused more while performing classification. A subset of obtained singular and token pairs for the tweet "Yay for getting pink eye again! #whyme" can be seen in Table 3 and 4 respectively.

3.2. Adding Ironic Tokens to Non-ironic Text

For non ironic text, we randomly sampled 100 tweets, labeled as non ironic, from the train and test datasets.

To add the ironic tokens into the non-ironic text we simply appended them to the end of the tweet. This was done for each token in the list applied to each tweet in the subset. We are aware that simply appending the token, or tokens, to the end of a tweet may generate something meaningless but ensuring that the newely formed tweet is semanticaly correct is out of the scope of this paper and we leave that for future work.

4. Results

4.1. Single Token

We analysed (X) tweets paired with 10 tokens. A sample of the top (N) and bottom (M) changes in non-ironic confidence scores can be seen in (ADD TABLE), with the full results available at: (LINK TO GITHUB).

The results show that the token with the most influence on the ironic sentiment was (X) and that token (Y) emphasizes the non ironic sentiment of a tweet. (ADD MORE COMMENTS ON THE RESULTS)

4.2. Token Pairs

We analysed 100 tweets paired with 200 pairs of tokens, producing a total of 20 000 new tweets. A sample of the top (N) and bottom (M) changes in non-ironic confidence scores can be seen in (ADD TABLE), with the full results available at: (LINK TO GITHUB).

The results show that the token pair with the most influence on the ironic sentiment was (X) and that token pair (Y) emphasizes the non ironic sentiment of a tweet. (ADD MORE COMMENTS ON THE RESULTS)

5. Future Work

Even though we have managed to change ironic sentiment by applying ironic tokens to tweets, we propose a more robust approach that takes into account that the semantics of a tweet remain intact after adding the token. Additionally, possible exploration of adding tokens at arbitrary points in the tweet, not only at the end.

As mentioned, a higher attention score doesn't necessarily mean that a token is more ironic, thus in future work, we could explore different metrics for obtaining tokens which we want to analyse.

6. Conclusion

Conclusion is the last enumerated section of the paper. It should not exceed half of a column and is typically split into 2–3 paragraphs. No new information should be presented in the conclusion; this section only summarizes and concludes the paper.

References

S. Abnar and W. Zuidema. 2020. Quantifying attention flow in transformers. *arXiv preprint arXiv:2005.00928*.

C. Van Hee, E. Lefever, and V. Hoste. 2018. SemEval-2018 task 3: Irony detection in English tweets. In *Proceedings of The 12th International Workshop on Semantic Evaluation*, pages 39–50, New Orleans, Louisiana, June. Association for Computational Linguistics.

Roger J. Kreuz. 2020. What makes something ironic? *The Conversation*, February.

Y. Liu, M. Ott, N. Goyal, J. Du, M. Joshi, D. Chen, O. Levy, M. Lewis, L. Zettlemoyer, and V. Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*.

S. Serrano and N. A. Smith. 2019. Is attention interpretable? *arXiv preprint arXiv:1906.03731*.

A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems*, 30.

Table 1: Most Ironic Tweets in the Training Set

Tweet	Confidence score
Yay for getting pink eye again! #whyme	1
@user yay!!!! It works #HateWhenThingsDontWorkRight	0.99
@user great Christmas.	0.99
Well done @user for making it possible to get emergency messages to a member of staff.	0.99
Isn't it great to sleep 5 hours and feel like a million bucks? #gettingold	0.99

Table 2: Most Ironic Tweets in the Test Set

Tweet	Confidence score
Just great when you're mobile bill arrives by text	1
OH and now the District line has major signal failures and delays FANTASTIC!!!	0.99
A wonderful day of starting work at 6am	0.99
Having to be up in four hours sounds great	0.99
@user @user @user @user @user nice to see the ambulance service is so important to OUR mps	0.99

Table 3: Single Tokens With Highest Attention

Token	Attention score
me	1
Y	0.37
pink	0.33
ay	0.29
why	0.14

Table 4: Token Pairs With Highest Attention

Token1	Token2	Attention score
ay	Y	0.6
#	Y	0.6
␣	Y	0.6
me	Y	0.6
again	Y	0.6